

1 E2017120501: Data Partitioning on Fragmented XML Data

1.1 Settings

The settings of this experiments are listed below.

Hardware There were 1 master (HaoDesk) and 16 workers (matsu-lab00 – matsu-lab15) used in the experiments listed in A.2.

Software The input XML data set was xmark600(see A.2), which was fragmented into 16 fragments with maxsize=4M. On each worker, there ran a BaseX server (Verion: 8.6.7). For each fragment, a BaseX database instance, namely ‘xmark600.16_4M’ was created in main memory on each worker. For the data partition, the number of partions(P) were 1, 2, 3, 4, 8. The intermediate database for holding the results of prefix query were named ‘xmark600.16_4M.tmp’.

XQuery Expressions Two XPath expressions XQ1 and XQ2 for prefix and suffix queries are listed in Section A.2.

1.2 Confirming Correctness

The number of hit nodes, the order of hit nodes and results size of the final output of queries have been checked and confirmed to be correct. All the successfully evaluated queries have the same number of hit nodes in original order. Some may be different in size, such as xm3a.dps. Its original result size was 922,270,281, but in this experiment, it is 883,253,777. This dramatical difference is caused by line-break. The previous experiments were done on Windows, while the current on Linux. Since BaseX using ‘\r\n’ on Windows, while ‘\n’ on Linux for line-break, there is one byte difference for each new line. We also found that the query has 6,502,751 hit nodes and there are six lines in every hit node, such as:

```
<bidder>
  <date>08/04/1999</date>

  <time>11:15:36</time>
  <personref person="person17793"/>
  <increase>7.50</increase>
</bidder>
```

We then have $883,253,777 + 6,502,751 * 6 - 922,270,281 = 2$. These two bytes are extra ‘\n’. So the sizes are the same.

Tab. 1: My caption

Key	Prefix (s)	Suffix (s)					Merge (s)	Original (s)	Results Size	
		P=1	P=2	P=3	P=4	P=8			hit nodes(M)	in MiB
xm1.dps								3122	5.5	57267
xm2.dps								0.01	0	0
xm3.dps	0.3	113	58.5	30.1	20.2	13.5	6	63.3	6.5	880
xm4.dps								101	0.6	1511
xm5.dps	0.3	243	67.5	48.9	37.8	31.8	10	74.9	35.9	944
xm6.dps								70	1.3	1289

1.3 Discussion on Queries

xm1.dps Since the results of xm1.dps are larger than the memory of HaoDesk. This query was not evaluated. This is a design choice about how to process the results. Based the previous evaluation, one possible way is to stored the unordered intermediate results in files and then concatenate these files.

xm2.dps Not tested for two reasons. Firstly, the prefix query of xm2.dps (as well as xm4.dps) cannot be processed by the current XQuery expression XQ1 (will be modified and tested soon). Secondly, there is no results as shown in experiment E20170401. One proposal is to make a minimal change to the query, such as change `category52` to `category324329`, which exists in XMark600.

xm3.dps The results of xm3.dps are shown in Table ?? (xm3.org = 63.32s)

xm4.dps Not tested for the same reason as the first one of xm2.dps.

xm5.dps The query xm5a.dps was the targe query in the previous design. However, due to the insufficient number of nodes in the results of prefix query to be allocated to 16 computers¹, I changed it to xm5b.dps (xm5c.dps was also evaluated, but it took 1480s (P=4), which is too long and thus ignored). The results of xm5b.dps is listed in Table ?? (xm5.dps = 75.05s).

xm6.dps Failed to evaluate them with data partitioning. The reason is the same as xm5a.dps (both `/site/regions` and `/site/regions/*[...]/item` were tested).

¹ Detailed explanation: An error message “database ‘xmark600_16.4M.tmp’ has no node with pre value 5.” was encountered when executing a suffix query on an intermediate database. The database had the content of `<root><part> ...</part></root>`, where there is only one `part` node. In a suffix query when $P = 2$, it is then to process “pre value 5” in XQ2, which refers to the second `part` node. However, since there is no second `part` node, the error occurred.

Tab. 2: Original XPath queries on XMark datasets.

key	query
xml.org	/site//*[name(.)="emailaddress" or name(.)="annotation" or name(.)="description"]
xm2.org	/site//incategory[./@category="category52"]/parent::item/@id
xm3.org	/site/open_auction/bidder[last()]
xm4.org	/site/regions/*/item[./location="United States" and ./quantity > and ./payment="Creditcard" and ./description and ./name]
xm5.org	/site/open_auctions/open_auction/bidder/increase
xm6.org	/site/regions/*[name(.)="africa" or name(.)="asia"]/item/description/parlist/listitem

1.4 Efficiency of Parsing Intermediate Results of Suffix Query

A important method in the implemenation that affects the whole performance is `baseex.PreValueReceiver.process(InputStream input)` used to parse the received results of suffix query, (i.e. PRE value + content). It takes the results of suffix query and returns an `QueryResultsPre` instance with a list of PRE values and a list of string content (of the same size). A simple experiment were done to evaluate the parsing speed. In the experiment, it took 465 ms to parse 52,757 KiB data with 704,430 nodes, i.e. it can process 100 MiB data per second, which basically reach the maximum network speed and thus should be sufficient for not being a bottleneck.

A Queries

A.1 XPath Quereis

The original XPath queries for XMark datasets are listed in Table 2. The data partitioning XPath queries of Table 2 is listed in Table 3.

A.2 XQuery Expressions

The two XQuery expressions are used for processing prefix query and suffix query for data partitioning strategy shown in Table A.2 and Table A.2 respectively. Note that `{eval_prefix}` will be first replaced by `{prefix}` or `db:open('{db}') {prefix}` depending on whether the query is optimized.

Tab. 3: Data partitioning XPath Queries on XMark datasets.

key	query
xml1.dps	pre = /site//*, suf = [name(.)="emailaddress" or name(.)="annotation" or name(.)="description"]
xm2.dps	pre = db:attribute("{db}", "category52"), suf = /parent::incategory/parent::item/@id
xm3.dps	pre = /site/open_auctions/open_auction, suf = /bidder[last()]
xm4.dps	pre = db:text("{db}", "Creditcard") suf = /parent::*:item[parent::*:parent::*:regions /parent::*:site/parent::document-node()] [(*:location = "United States")][0.0 < *:quantity] [*:description][*:name]
xm5a.dps	pre = /site/open_auctions, suf = /open_auction/bidder/increase
xm5b.dps	pre = /site/open_auctions/open_auction, suf = /bidder/increase
xm5c.dps	pre = /site/open_auctions/open_auction/bidder, suf = /increase
xm6.dps	pre = /site/regions, suf = /self::*[name(.)="africa" or name(.)="asia"]/item /description/parlist/listitem

Tab. 4: XQuery Expression for XQ1 prefix part

```
// XQ12: for prefix query
let $d := array { {eval_prefix} !
    db:node-pre(.) } return
for $i in 0 to {P} - 1 return
let $q := array:size($d) idiv {P} return
let $r := array:size($d) mod {P} return
let $part_length := if ($i < $r) then $q + 1
else $q return
let $part_begin := if ($i <= $r) then ($q + 1) * $i
else $q * $i + $r return
insert node element part {
array:subarray($d, $part_begin + 1, $part_length)
} as last into db:open('{tmpdb}')/root
```

Tab. 5: XQuery Expression XQ2 for suffix part

```

// XQ2: for suffix query
declare option output:method '$mode';
declare option output:item-separator '[';

for $pre in db:open('{tmpdb}')/root return
let $node := db:open('{db}') { suffix }
return (db:node-pre($node), $node)

let $part_pre := {p}*2 + 1 return
for $pre in ft:tokenize(db:open-pre('{tmpdb}', $part_pre)) return
for $node in db:open-pre('{db}', xs:integer($pre)) { suffix }
return (db:node-pre($node), $node)

```